


Practical Question 6: Configure Auto Scaling and Load Balancing

Objective: Set up an auto-scaling group and load balancer.

Steps:

a) Create an Auto Scaling Group and define a launch template.



The screenshot shows the AWS Management Console interface for creating a new launch template. The breadcrumb navigation at the top indicates the path: **EC2** > **Launch templates** > **Create launch template**. The main heading is **Create launch template**, followed by a descriptive paragraph: "Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions." The form is divided into several sections: 1. **Launch template name and description**: Includes a text input for the name (containing "myTemplate") with a note that it must be unique and 128 characters or less, and a text input for the description (containing "A prod webserver for MyApp") with a note that it can be up to 255 characters. 2. **Auto Scaling guidance**: A section with an "Info" link and a checkbox labeled "Provide guidance to help me set up a template that I can use with EC2 Auto Scaling", which is currently unchecked. 3. **Template tags**: A section with a right-pointing triangle icon. 4. **Source template**: A section with a right-pointing triangle icon. The bottom of the page shows the start of the **Launch template contents** section.

Configure- AMI: Amazon Linux / Ubuntu

☰ [EC2](#) > [Launch templates](#) > Create launch template

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI contains the operating system, application server, and applications for your instance. If you don't see a suitable AMI below, use the search field or choose **Browse more AMIs**.

🔍 Search our full catalog including 1000s of application and OS images

Recents **Quick Start**

Don't include in launch template

Amazon Linux
aws

macOS
Mac

Ubuntu
ubuntu

Windows
Microsoft

Red Hat
Red Hat


Browse more AMIs
Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2023 kernel-6.1 AMI
ami-0a0823e4ea064404d (64-bit (x86), uefi-preferred) / ami-069ac965b8beaa442 (64-bit (Arm), uefi)
Virtualization: hvm ENA enabled: true Root device type: ebs

Free tier eligible ▼

Instance type: t3.micro (free tier)
Key pair: Select or create.

☰ [EC2](#) > [Launch templates](#) > Create launch template

▼ Instance type [Info](#) | [Get advice](#) Advanced

Instance type

t3.micro
Family: t3 2 vCPU 1 GiB Memory Current generation: true
On-Demand Linux base pricing: 0.0108 USD per Hour
On-Demand Windows base pricing: 0.02 USD per Hour
On-Demand Ubuntu Pro base pricing: 0.0143 USD per Hour
On-Demand RHEL base pricing: 0.0396 USD per Hour
On-Demand SUSE base pricing: 0.0108 USD per Hour

Free tier eligible ▼

☒ All generations

[Compare instance types](#)

Additional costs apply for AMIs with pre-installed software

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name

ec2 ▼

 [Create new key pair](#)

Security group: Allow **SSH (22)**
Allow **HTTP (80)**

Info

- ☐ Select existing security group

MySecurityGroup

Info

Allow SSH (22) Allow HTTP (80)

Info


```
vpc-0cf5a8304129d39ac
172.31.0.0/16
```

(default) ▼




Add security group rule

Add User Data.




Click **Create template.**



Step B: Create Auto Scaling Group.
Using Amazon EC2 Auto Scaling:

1. Go to EC2 → Auto Scaling Groups.
2. Click Create Auto Scaling Group.
3. Select your Launch Template.



Configure: VPC: Default or custom.
Subnets: Select at least 2 (for high availability).

Network [Info](#)

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC
Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-0cf5a8304129d39ac
172.31.0.0/16 Default

[Create a VPC](#)

Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

eun1-az1 (eu-north-1a) | subnet-02257f987dd57ab77
172.31.16.0/20 Default

eun1-az2 (eu-north-1b) | subnet-0515a5cdbc03539cb
172.31.32.0/20 Default

[Create a subnet](#)

Availability Zone distribution - new
Auto Scaling automatically balances instances across Availability Zones. If launch failures occur in a zone, select a strategy.

☒ **Balanced best effort**
If launches fail in one Availability Zone, Auto Scaling will attempt to launch in another healthy Availability Zone.

☐ **Balanced only**
If launches fail in one Availability Zone, Auto Scaling will continue to attempt to launch in the unhealthy Availability Zone to preserve balanced distribution.

© 2026, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Set group size:

Desired: 1

Min: 1

Max: 3 (or more)

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1
● Choose launch template

Step 2
● Choose instance launch options

Step 3 - optional
● Integrate with other services

Step 4 - optional
● **Configure group size and scaling**

Step 5 - optional
● Add notifications

Step 6 - optional
● Add tags

Step 7
● Review

Configure group size and scaling - optional [Info](#)

Define your group's desired capacity and scaling limits. You can optionally add automatic scaling to adjust the size of your group.

Group size [Info](#)
Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.

Desired capacity type
Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.

Units (number of instances)

Desired capacity
Specify your group size.

1

Scaling [Info](#)
You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits
Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity
1
Equal or less than desired capacity

Max desired capacity
3
Equal or greater than desired

CloudShell Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Step C: Configure Application Load Balancer.

Using Application Load Balancer:

1. Choose **Attach to new Load Balancer**

2. Select:

- Type: **Application Load Balancer (ALB)**
- Scheme: Internet-facing

Create Application Load Balancer [Info](#)

The Application Load Balancer distributes incoming HTTP and HTTPS traffic across multiple targets such as Amazon EC2 instances, microservices, and containers, based on request attributes. When the load balancer receives a connection request, it evaluates the listener rules in priority order to determine which rule to apply, and if applicable, it selects a target from the target group for the rule action.

► **How Application Load Balancers work**

Basic configuration

Load balancer name
Name must be unique within your AWS account and can't be changed after the load balancer is created.

MyLoadBalancer

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme [Info](#)
Scheme can't be changed after the load balancer is created.

☒ **Internet-facing**

- Serves internet-facing traffic.
- Has public IP addresses.
- DNS name resolves to public IPs.
- Requires a public subnet.

☐ **Internal**

- Serves internal traffic.
- Has private IP addresses.
- DNS name resolves to private IPs.
- Compatible with the IPv4 and Dualstack IP address types.

Load balancer IP address type [Info](#)
Select the front-end IP address type to assign to the load balancer. The VPC and subnets mapped to this load balancer must include the selected IP address types. Public IPv4 addresses have an additional cost.

CloudShell Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Configure: Listener: HTTP (port 80).

Availability zones: Same as ASG.

Network mapping [Info](#)

The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.

VPC [Info](#)
The load balancer will exist and scale within the selected VPC. The selected VPC is also where the load balancer targets must be hosted unless routing to Lambda or on-premises targets, or if using VPC peering. To confirm the VPC for your targets, view [target groups](#).

vpc-0cf5a8304129d39ac (default) [Create VPC](#)

IP pools [Info](#)
You can optionally choose to configure an IPAM pool as the preferred source for your load balancers IP addresses. Create or view Pools in the [Amazon VPC IP Address Manager console](#).

☐ **Use IPAM pool for public IPv4 addresses**
The IPAM pool you choose will be the preferred source of public IPv4 addresses. If the pool is depleted IPv4 addresses will be assigned by AWS.

Availability Zones and subnets [Info](#)
Select at least two Availability Zones and a subnet for each zone. A load balancer node will be placed in each selected zone and will automatically scale in response to traffic. The load balancer routes traffic to targets in the selected Availability Zones only.

☒ **eu-north-1a (eu-n1-az1)**

Subnet
Only CIDR blocks corresponding to the load balancer IP address type are used. At least 8 available IP addresses are required for your load balancer to scale efficiently. Selected subnets must have Network ACL inbound and outbound ALLOW rules, and internet gateway (IGW) routes configured to allow traffic.

subnet-02257f987dd57ab77
172.31.16.0/20
Network ACL ALLOW rules: Inbound (1) Outbound (1) | Route table: IGW (-/)

☒ **eu-north-1b (eu-n1-az2)**

Subnet
Only CIDR blocks corresponding to the load balancer IP address type are used. At least 8 available IP addresses are required for your load balancer to scale efficiently. Selected subnets must have Network ACL inbound and outbound ALLOW rules, and internet gateway (IGW) routes configured to allow traffic.

subnet-0515a5cde03539cb
172.31.32.0/20
Network ACL ALLOW rules: Inbound (1) Outbound (1) | Route table: IGW (-/)

CloudShell Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Create **Target Group**: Target type: Instances

aws

Search

[Alt+S]

Europe (Stockholm)

Ayush Kumar (031363453628)

Ayush Kumar

EC2 > Target groups > Create target group

Step 1: Create target group

Step 2 - recommended: Register targets

Step 3: Review and create

Create target group

A target group can be made up of one or more targets. Your load balancer routes requests to the targets in a target group and performs health checks on the targets.

Settings - immutable
 Choose a target type and the load balancer and listener will route traffic to your target. These settings can't be modified after target group creation."

Target type
 Indicate what resource type you want to target. Only the selected resource type can be registered to this target group.

☒ **Instances**
 Supports load balancing to instances in a VPC. Integrate with Auto Scaling Groups or ECS services for automatic management.
 Suitable for: ALB NLB GWLB

☐ **IP addresses**
 Supports load balancing to VPC and on-premises resources. Facilitates routing to IP addresses and network interfaces on the same instance. Supports IPv6 targets.
 Suitable for: ALB NLB GWLB

☐ **Lambda function**
 Supports load balancing to a single Lambda function. ALB required as traffic source.
 Suitable for: ALB

☐ **Application Load Balancer**
 Allows use of static IP addresses and PrivateLink with an Application Load Balancer. NLB required as traffic source.
 Suitable for: NLB

Target group name
 Name must be unique per Region per AWS account.

Accepts: a-z, A-Z, 0-9, and hyphen (-). Can't begin or end with hyphen. 1-32 total characters; Count: 8/32

Protocol
 Protocol for communication between the load balancer and targets.

Port
 Port number where targets receive traffic. Can be overridden for individual targets during registration.

CloudShell Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Health check path: /

Health checks

The associated load balancer periodically sends requests, per the settings below, to the registered targets to test their status.

Health check protocol

Health check path
 Use the default path of "/" to perform health checks on the root, or specify a custom path if preferred.

Up to 1024 characters allowed.

► Advanced health check settings

aws

Search

[Alt+S]

Europe (Stockholm)

Ayush Kumar (031363453628)

Ayush Kumar

EC2 > Load balancers > Create Application Load Balancer

Listeners and routing Info

A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.

▼ Listener HTTP:80

Remove

Protocol

HTTP

Port

80

1-65535

Default action Info
 The default action is used if no other rules apply. Choose the default action for traffic on this listener.

Routing action

☒ Forward to target groups
 ☐ Redirect to URL
 ☐ Return fixed response

Forward to target group Info
 Choose a target group and specify routing weight or [create target group](#).

Target group	Weight	Percent
MyTarget Target type: Instance, IPv4 Target stickiness: Off	1 0-999	100%


+ Add target group
 You can add up to 4 more target groups.

Target group stickiness Info
 Enables the load balancer to bind a user's session to a specific target group. To use stickiness the client must support cookies. If you want to bind a user's session to a specific target, turn on the Target Group attribute Stickiness.

☐ Turn on target group stickiness

CloudShell Feedback Console Mobile App


© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences



Step D: Configure Scaling Policies.

In Auto Scaling Group → **Automatic Scaling**.

Click **Add policy**.



Choose: **Target tracking policy**.

- Metric: **Average CPU Utilization**
- Target value: **70%**

Save policy.

aws

Search [Alt+S]

Europe (Stockholm)

Ayush Kumar (031363453628)

Ayush Kumar

EC2 > Auto Scaling groups > ScalingGroup > Dynamic scaling policy

Create dynamic scaling policy

Policy type

Target tracking scaling

Scaling policy name

Target Tracking Policy

Metric type [Info](#)

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

Target value

70

Instance warmup [Info](#)

300 seconds

☐ Disable scale in to create only a scale-out policy

Cancel Create

CloudShell Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step E: Test Auto Scaling (Simulate Traffic)

Method 1: Manual Load Test

SSH into instance and run: `sudo yum install -y stress`

`stress --cpu 2 --timeout 300`

aws

Search [Alt+S] Ask Amazon Q

Europe (Stockholm)

Ayush Kumar (031363453628)

Ayush Kumar

EC2 > Instances > i-058acefe22b4e8f92 > Connect to instance

Connect [Info](#)

Connect to an Instance using the browser-based client.

Instance ID i-058acefe22b4e8f92	VPC ID vpc-0cf5a8304129d39ac	Security groups sg-016ab1063eaea7589 (MySecurityGroup)	IAM role -
---	--	--	----------------------

EC2 Instance Connect SSM Session Manager SSH client EC2 serial console

Instance ID
i-058acefe22b4e8f92

Connection type

☒ Connect using a Public IP
Connect using a public IPv4 or IPv6 address

☐ Connect using a Private IP
Connect using a private IP address and a VPC endpoint

Public IPv4 address
16.170.205.74

☐ IPv6 address
-

Username
Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

ec2-user

Note: In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

CloudShell Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

```
aws
Search [Alt+S]
Europe (Stockholm) Ayush Kumar (031363453628) Ayush Kumar

Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

Last login: Sun Apr 26 05:14:21 2026 from 13.48.4.202
[ec2-user@ip-172-31-46-170 ~]$ sudo yum install -y stress
Last metadata expiration check: 14:35:25 ago on Sat Apr 25 14:41:44 2026.
Package stress-1.0.7-2.amzn2023.0.1.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@ip-172-31-46-170 ~]$ stress --cpu 2 --timeout 300
stress: info: [63966] dispatching hogs: 2 cpu, 0 io, 0 vm, 0 hdd
```

Step F: Verify Scaling


1. Go to **Auto Scaling Group** → **Activity**
2. Check:
 - New instances launching
3. Go to **EC2 Instances**
 - See instances increasing/decreasing.

The screenshot displays the AWS Management Console interface for an Auto Scaling Group. The 'Activity' tab is selected, showing a table of activity history. The table has columns for Status, Description, Cause, Start time, and End time. A single activity is listed with a 'Successful' status, describing the launch of a new EC2 instance. The console also includes navigation tabs for Details, Integrations, Automatic scaling, Instance management, Instance refresh, Activity, Monitoring, and Tags. At the bottom, there are links for CloudShell, Feedback, Console Mobile App, and footer information including copyright and privacy links.

Status	Description	Cause	Start time	End time
Successful	Launching a new EC2 instance: i-058acefe22b4e8f92	At 2026-04-25T14:38:16Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 1. At 2026-04-25T14:41:16Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 1.	2026 April 25, 08:11:17 PM +05:30	2026 April 25, 08:11:23 PM +05:30

4. Check Load Balancer:

- Targets becoming healthy.




Practical Question 7: Monitor Resources Using AWS CloudWatch

Objective: Use CloudWatch to monitor AWS resources.

Step A: Set Up CloudWatch Metrics (EC2 CPU Monitoring)


1. Login to AWS Management Console.
2. Go to **EC2 Dashboard** → *Instances*.
3. Select your running instance.
4. Click the **Monitoring** tab.
5. View metrics like:
 - CPU Utilization.
 - Network In/Out.
 - Disk Read/Write.



Step B: Create a CloudWatch Alarm

Go to **CloudWatch Dashboard**.

Click **Alarms** → **Create Alarm**.



Click **Select metric**.

Select metric

Untitled graph

1h3h12h1d3d1wCustomUTC timezoneLine

1

Your CloudWatch graph is empty.
Select some metrics to appear here.

0.5

0

15:1515:3015:4516:0016:1516:3016:4517:0017:1517:3017:4518:00

Browse (595)Multi source queryGraphed metricsOptionsSourceAdd mathAdd query

Alarm recommendationsGraph with SQLGraph search

Search for any metric, dimension, resource id or account id

ApplicationELB10

EBS45

EC286

Logs4

S32

States8

Usage440

CancelSelect a single metric to continue

Choose:

EC2 → Per-Instance Metrics.

Select **CPUUtilization** for your instance

Select metric

CPUUtilization

1h3h12h1d3d1wCustomUTC timezoneLine

1

No data available.
Try adjusting the dashboard time range.

0.5

0

15:1515:3015:4516:0016:1516:3016:4517:0017:1517:3017:4518:00

CPUUtilizationBrowse (66)Multi source queryGraphed metrics (1)OptionsSourceAdd mathAdd query

☐

i-01fec2a892a80a246

MetadataNoToken

No alarms

☐

i-01fec2a892a80a246

StatusCheckFailed_AttachedEBS

No alarms

☐

i-01fec2a892a80a246

InstanceEBSIOPEXceededCheck

No alarms

☐

i-01fec2a892a80a246

EBSReadBytes

No alarms

☐

i-01fec2a892a80a246

StatusCheckFailed_Instance

No alarms

☒

i-01fec2a892a80a246

CPUUtilization

No alarms

☐

i-01fec2a892a80a246

StatusCheckFailed

No alarms

☐

i-0341ca599c08b0099

EBSReadBytes

No alarms

CancelSelect metric

Click Select Metric.

aws

Search

[Alt+S]

Europe (Stockholm)

Ayush Kumar (031363453628)

Ayush Kumar

CloudWatch

Alarms

Create alarm

Step 1

Specify metric and conditions

Step 2

Configure actions

Step 3

Add alarm details

Step 4

Preview and create

Specify metric and conditions

Alarm recommendations

View details

Metric

Edit

Graph

This alarm will trigger when the blue line goes above the red line for 1 datapoints within 5 minutes.

1

No data available.
Try adjusting the dashboard time range.

0.5

0

15:3016:0016:3017:0017:3018:00

CPUUtilization

Namespace

AWS/EC2

Metric name

CPUUtilization

InstanceId

i-01fec2a892a80a246

Instance name

No name specified

Statistic

Average

Period

5 minutes

CloudShell

Feedback

Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

Configure Alarm:

Set:

Threshold type: Static.

Condition: **Greater than 70%.**

Click **Next**.

Conditions

Threshold type

Static

Use a value as a threshold

Anomaly detection

Use a band as a threshold

Whenever CPUUtilization is...

Define the alarm condition.

Greater

> threshold

Greater/Equal

>= threshold

Lower/Equal

<= threshold

Lower

< threshold

than...

Define the threshold value.

70

Must be a number.

Additional configuration

Cancel

Next

Step C: Configure SNS Topic (Email Notification)


1. In alarm setup → choose **Create new topic**.
2. Enter: Topic name (e.g., *CPU-Alert-Topic*) and Email address (your email).
3. Click **Create topic**.

Check your email inbox → **Confirm subscription**.

The screenshot shows the AWS CloudWatch 'Create alarm' console, specifically the 'Configure actions' step. The 'Notification' section is highlighted, showing three alarm state triggers: 'In alarm' (selected), 'OK', and 'Insufficient data'. Below this, the 'Send a notification to the following SNS topic' section shows 'Create new topic' selected. The 'Create a new topic...' section has 'CPU_alert_topic' entered. The 'Email endpoints that will receive the notification...' section has 'ayushkumar1502@gmail.com' entered. Buttons for 'Create topic' and 'Add notification' are at the bottom.

Step D: Complete Alarm Setup

1. Give alarm name (e.g., *High-CPU-Alarm*).
2. Click **Create Alarm**.




Step E: Test the Setup (Simulate High CPU Usage)

Connect to EC2 instance via SSH and run:

```
sudo yum install -y stress
```

```
stress --cpu 2 --timeout 300
```

This will increase CPU usage.



```
aws
Search [Alt+S]
Europe (Stockholm) Ayush Kumar (031363453628) Ayush Kumar

Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

Last login: Sun Apr 26 05:16:43 2026 from 13.48.4.202
[ec2-user@ip-172-31-46-170 ~]$ sudo yum install -y stress
Last metadata expiration check: 14:59:12 ago on Sat Apr 25 14:41:44 2026.
Package stress-1.0.7-2.amzn2023.0.1.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@ip-172-31-46-170 ~]$ stress --cpu 2 --timeout 500
stress: info: [65342] dispatching hogs: 2 cpu, 0 io, 0 vm, 0 hdd
```

i-058acefe22b4e8f92
PublicIPs: 16.170.205.74 PrivateIPs: 172.31.46.170

CloudShell Feedback Console Mobile App © 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step F: Verify Alarm and Notification

1. Go to **CloudWatch** → **Alarms**.
2. Observe: Alarm state changes to **In Alarm**.

aws

Search

[Alt+S]

Europe (Stockholm)

Ayush Kumar (031363453628)

Ayush Kumar

CloudWatch

Alarms

Alarms

Mute rules - new

Alarms (3)

Clear selection

Create composite alarm

Actions

Create alarm

Filter alarms

Quick filters

< 1 >

	Name	State	Actions	Actions muted - new	Last state update (UTC)	Conditions
☐	TargetTracking-ScalingGroup-AlarmLow-b7782bf6-853d-4769-a194-7bc032dc0336	⚠ In alarm	✅ Actions enabled	-	2026-04-26 05:31:22	CPUUtilization < 49 for 15 data minutes
☐	CPU_Alarm	⌚ Insufficient data	✅ Actions enabled	-	2026-04-25 18:22:15	CPUUtilization > 70 for 1 datap minutes
☐	TargetTracking-ScalingGroup-AlarmHigh-b0f4e95d-ea38-4d1f-a10f-03a4b67ed957	✅ OK	✅ Actions enabled	-	2026-04-25 15:22:51	CPUUtilization > 70 for 3 datap minutes

CloudShell

Feedback

Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences